

210 Systems Programming

File Types, Access Rights, Globbing, Text Editors, and I/O Redirection

Fall 2025

Week 2

Overview

- File types and access rights
- Globbing
- Text editors: vim, emacs, nano.
- I/O redirection

File Types

- We will learn about the file system in detail later. For now, you should know that there are:
 - *Text files*: they can be edited with text editors and displayed with commands such as `less`. Examples: source code, configuration files.
 - *Binary files*: they have different formats and cannot be edited/displayed directly. Examples: executable programs, `.docx` files, `.jpeg` files.
 - *Directories*
 - *Symbolic links*: shortcuts to files or directories (more about them later)
- By just looking at the file name, it is not possible to distinguish files/directories, or text/binary files.
 - File extensions are for informative purposes only. There can be files with no extensions or directories with extensions. (Example: try `"mkdir a.txt"`)

Multi-user Operating Systems

- Multi-user OSs like Unix/Linux should have mechanisms to protect each user's privacy and security
 - Where are each user's files located?
 - Who can access them? (What is "access"?)

The Unix Approach

- Organize *users* into *groups*
 - Each user has a unique user id and can be a member of multiple groups.
 - Use the `id` command to get information on a user:

```
id tolgacan
uid=430462(tolgacan) gid=430462(tolgacan)
groups=430462(tolgacan),869(mioaccess),882(wendian)
```
- Create separate read/write/execute permissions for the *owner* (u), the *group* (g), and all the *others* (o)
 - We can use the `chmod` command to change access permissions of files and directories

Permission Syntax

- Three permission types: read, write, execute (rwx) for three classes: owner, group, others (ugo) \Rightarrow 9 permission flags:
 - $r_u w_u x_u r_g w_g x_g r_o w_o x_o$
- Examples:
 - user has rwx, group has r, others have r \Rightarrow 111 100 100 or `rwxr--r--`
 - user has rw, group has rw, others have none \Rightarrow 110 110 000 or `rw-rw----`

Querying File Permissions

- `ls -l` will show file permissions in the first column as a 10-character string.
 - The first character shows us whether it is a directory, file, or a symbolic link.
 - `d` for directories, `-` for files, and `l` for symbolic links.
- Example: `-rwxr-xr--`
 - Owner has all permissions
 - Users in the group have read and execute permissions
 - Others have read permission

Changing permissions with `chmod`

- Can use the *symbolic* or the *numeric* mode
 - The format of a symbolic mode is `[ugoa...][[-+=[perms...]]...]`, where `perms` is either zero or more letters from the set `rwXst`, or a single letter from the set `ugo`. Multiple symbolic modes can be given, separated by commas.
 - Examples:
 - `chmod +x foo`
 - `chmod u+rw,g+x,o-rwx foo`
 - The numeric mode use octal digits (0-7) to represent the 3-bit permissions for each `ugo` category.
 - Examples:
 - `chmod 754 foo`
 - `chmod 600 foo`

Execute permission on files vs on directories

- Execute permission on files means the right to execute them like commands if they are programs or shell scripts.
 - You should not give the execute permission to non-program/non-script files. If you do, the content will be treated as a sequence of commands.
- For directories, execute permission allows you to `cd` into that directory.

Practice

What are the permissions of each of the following files after each command is run? Assume initial permissions of `rw-rw-rw-`.

1 `chmod u=rwx,g=r,o= file2`

2 `chmod 441 foo`

3 `chmod u+x hello`

4 `chmod u+rwx,g-x world`

5 `chmod 643 bar`

Wildcards and Globbing

- Bash supports a system of pattern matching for referring to a set of files called *globbing*
- Although there are similarities, globbing does not use the standard Regular Expression syntax. Instead, globbing recognizes and expands wild cards, such as:
 - * and ?,
 - character lists in square brackets,
 - and certain other special characters (such as ^ for negating the sense of a match)
- Limitations:
 - Strings containing * will not match filenames that start with a dot.
Example: .bashrc.
 - The ? has a different meaning in globbing than as part of an RE.

Rules of Globbing

<code>*</code>	Matches 0 or more characters
<code>?</code>	Matches precisely one character
<code>[...]</code>	Matches specified characters in a set or in ranges
<code>[^...]</code>	Negate the set defined in square brackets
<code>[!...]</code>	Negate the set defined in square brackets
<code>{p1,p2,...}</code>	Match/expand each globbing pattern one by one

Note: In most shell implementations, `^` is used as the range negation character. However, POSIX specifies `!` for this role, and therefore `!` is the standard choice.

Examples

<code>*</code>	Matches any string, of any length
<code>foo*</code>	Matches any string beginning with foo
<code>*x*</code>	Matches any string containing an x (beginning, middle or end)
<code>*.tar.gz</code>	Matches any string ending with .tar.gz
<code>*.[ch]</code>	Matches any string ending with .c or .h
<code>foo?</code>	Matches foot or foo\$ but not fools
<code>[!ab]*</code>	Matches any string that doesn't start with a or b
<code>{b*,c*,*est*}</code>	Matches strings that start with b, start with c, and contain est in them.

When does globbing occur?

- Bash performs filename expansion, i.e., globbing, on unquoted command-line arguments.
- In other words, in commands such as `ls`, `mv`, and `cp`, expansion is performed to convert, for example,
 - `cp *.txt dir to`
 - `cp a.txt b.txt cc.txt dir`

More on globbing

- Bash (and Korn Shell) offers extended globs, which have the expressive power of regular expressions.
- Korn shell enables these by default; in Bash, you must run the following command to enable it:
 - `shopt -s extglob`
- Additional resources on globbing:
 - Globbing reference on the Linux Documentation Project - TLDP
 - The glob page on Greg's Wiki

Text editors

- We create text content, such as README files, code, Makefiles, to-do lists, and configuration files, usually using tools with GUIs such as VS Code, notepad++.
- Discussion:
 - What are text content do you create?
 - What other tools do you use?

Command line text editors

- What if you access a remote machine via a text terminal? Would these tools will be available to you then?
- Text editors that run from within a text based terminal are useful tools to create and modify such content.
- They also offer other advantages like speed and additional functionality.
- However, their learning curve can be steeper than tools with GUIs.

Vim



- Vim is a free and open-source, screen-based text editor program.
- It is an improved clone of Bill Joy's vi.
- Vim's author, Bram Moolenaar, derived Vim from a port of the Stevie editor for Amiga and released a version to the public in 1991.

Vim



- Vim is a modal editor
 - Save keypresses by switching modes instead of using complex shortcuts
- Normal mode
 - Default mode. Used for editor commands
- Visual mode
 - Similar to normal, but used to select areas of text

Vim



- Other modes:
 - Insert mode
 - Allows the user to type text into the file
 - Command mode
 - Shows a single line at the bottom for normal commands (save, quit, etc.)
- You can learn Vim by using the command line program `vimtutor`

Vim Pros and Cons



■ Pros:

- Installed on most systems (even on macOS)
- Very fast and lightweight
- Easily perform complex edits
- Shortcuts for everything, macros, registers, repetition, etc.

■ Cons:

- Steep learning curve

Emacs



- GNU Emacs is a free software text editor.
- It was created by GNU Project founder Richard Stallman, based on the Emacs editor developed for Unix operating systems.
- Its tag line is “the extensible self-documenting text editor.”

Emacs



- One of the oldest free software still in active development
- Uses sequences of shortcuts for operations
- Highly extensible and customizable
- You can learn emacs by launching it with `emacs` and then pressing `ctrl + h` and then `t`

Emacs Pros



■ Pros:

- Highly extensible and customizable
- Powerful (arguably more than any other editor, period)
- Mature integration with other free software tools
- You can use it for everything, even browsing the web!
- Choose Emacs if you want a complete development environment akin to VSCode, CLion, etc.

Emacs Cons



- Cons:
 - Bad ergonomics!
 - Emacs pinky syndrome
 - Recommendation: Use “Evil Mode” extension for Vim shortcuts
 - Not available everywhere. If you don’t have superuser privileges, you won’t be able to install

nano

```

001
iLE880j. :j0888880j:
.LGite8800.f06jjLE880E:
1E :888Et. .G888E:
;1 E888. 8888.
D888. :8888:
D888. :8888:
D888. :8888:
D888. :8888:
S88W. :8888:
W88W. :8888:
W88W. :8888:
DGGD: :8888:
:8888:
:W888:
E888i
tW888

```

- GNU nano is a text editor for Unix-like computing systems or operating environments using a command line interface.
- It emulates the Pico text editor, part of the Pine email client, and also provides additional functionality.
- Unlike Pico, nano is licensed under the GNU General Public License.

nano

```

      .L800j. :j000000j:
.LGitE000D.f06jjjL8000E;
.iE :E000E. ,G0000.
;i :E000. ,0000.
   D000. :0000;
   D000. :0000;
   D000. :0000;
   D000. :0000;
   000W. :0000;
   W00W. :0000;
   DGGD. :0000;
           :W000;
           :0000;
           E000;
           tW000;

```

- GNU nano is keyboard-oriented, controlled with control keys.
- GNU nano puts a two-line "shortcut bar" at the bottom of the screen, listing many of the commands available in the current context.
- `ctrl + g` gets you to the help screen.

nano Pros and Cons

```

      1111
4LE880j. :j088800j:
, LGitE880D. f06jjjLE880E;
1E :LE888ET. , G8888.
; i E888, , 8888,
D888, : 8888:
D888, : 8888:
D888, : 8888:
D888, : 8888:
888W, : 8888:
W88W, : 8888:
W88W, : 8888:
DGGD: : 8888:
: W888:
E888:
tW888

```

- Pros:

- Easier to learn
- GNU nano can also use pointing devices, such as a mouse, to activate functions that are on the shortcut bar, as well as position the cursor.

- Cons:

- Not as powerful as Emacs or Vim

Which editor to use?



```

      111
      1LE880j. :j088880j:
      ,LGtE880D.f06jjjLE888E;
      1E :8888E1. ,G8888.
      ;1 :8888. ,8888.
      D888. :8888:
      0888. :8888:
      D888. :8888:
      D888. :8888:
      888W. :8888:
      W88W. :8888:
      W88W. :8888:
      DGGD: :8888:
      :8888:
      :W888:
      E8881
      tW88D

```

- “See the Editor Wars!” at:
 - <https://www.linuxtrainingacademy.com/nano-emacs-vim/>

I/O Streams

- There are three standard input out streams recognized by operating systems and programming languages:
 - Standard Input: `stdin`
 - Standard Output: `stdout`
 - Standard Error: `stderr`
- You should be familiar with these from CSCI 200.
 - `cin`, `cout`, `cerr`
- These streams provide basic user input and program output for C/C++, Java, Python programs and Linux commands.

Standard I/O Default Values

The default values for standard input, standard output, and standard error are:

- Standard input: The terminal/console keyboard
- Standard output: The terminal/console display screen
- Standard error: The terminal/console display screen

How do these scale for large inputs and outputs?

Example scenario: testing a program

- You want to see whether your program has any bugs by testing it on a number of test cases.

Example scenario: testing a program

- You want to see whether your program has any bugs by testing it on a number of test cases.
- How many test cases would you test it against?

Example scenario: testing a program

- You want to see whether your program has any bugs by testing it on a number of test cases.
- How many test cases would you test it against?
- 10? 100? 1000?

Example scenario: testing a program

- You want to see whether your program has any bugs by testing it on a number of test cases.
- How many test cases would you test it against?
- 10? 100? 1000?
- How long would it take?

Example scenario: testing a program

- You want to see whether your program has any bugs by testing it on a number of test cases.
- How many test cases would you test it against?
- 10? 100? 1000?
- How long would it take?
- Resolve a bug and retest?

Example scenario: testing a program

- You want to see whether your program has any bugs by testing it on a number of test cases.
- How many test cases would you test it against?
- 10? 100? 1000?
- How long would it take?
- Resolve a bug and retest?
- How would you automate the testing process?

Automating Testing

- Possible solution:
 - Read the input from files and write the output to files

Automating Testing

- Possible solution:
 - Read the input from files and write the output to files
 - Would you need to change your program to be able to process files?

Automating Testing

- Possible solution:
 - Read the input from files and write the output to files
 - Would you need to change your program to be able to process files?
 - **No need.** There is an easier way: **I/O redirection**

I/O Redirection

- You can *redirect* standard input, standard output, and standard error to files using I/O redirection

I/O Redirection

- You can *redirect* standard input, standard output, and standard error to files using I/O redirection
- `myProg < input` tells the program that the standard input is no longer the keyboard. It is the file `input`.

I/O Redirection

- You can *redirect* standard input, standard output, and standard error to files using I/O redirection
- `myProg < input` tells the program that the standard input is no longer the keyboard. It is the file `input`.
 - The characters in the file will be processed as if the user typed them with a keyboard.

I/O Redirection

- You can *redirect* standard input, standard output, and standard error to files using I/O redirection
- `myProg < input` tells the program that the standard input is no longer the keyboard. It is the file `input`.
 - The characters in the file will be processed as if the user typed them with a keyboard.
- `myProg > output` tells the program that the standard output is no longer the screen. It is the file `output`.

I/O Redirection

- You can *redirect* standard input, standard output, and standard error to files using I/O redirection
- `myProg < input` tells the program that the standard input is no longer the keyboard. It is the file `input`.
 - The characters in the file will be processed as if the user typed them with a keyboard.
- `myProg > output` tells the program that the standard output is no longer the screen. It is the file `output`.
 - The output of the program will be directly written to the file `output`.

I/O Redirection

- You can *redirect* standard input, standard output, and standard error to files using I/O redirection
- `myProg < input` tells the program that the standard input is no longer the keyboard. It is the file `input`.
 - The characters in the file will be processed as if the user typed them with a keyboard.
- `myProg > output` tells the program that the standard output is no longer the screen. It is the file `output`.
 - The output of the program will be directly written to the file `output`.
- Can do both simultaneously: `myProg < input > output`

Back to the testing problem

- How would you use I/O redirection to speed up the testing process?

Back to the testing problem

- How would you use I/O redirection to speed up the testing process?
 - Write a test generator to generate input files (Use output redirection).

Back to the testing problem

- How would you use I/O redirection to speed up the testing process?
 - Write a test generator to generate input files (Use output redirection).
 - Run your program with I/O redirection producing output files.

Back to the testing problem

- How would you use I/O redirection to speed up the testing process?
 - Write a test generator to generate input files (Use output redirection).
 - Run your program with I/O redirection producing output files.
 - Write a program to compare the generated outputs against a set of expected outputs (Can use existing Linux utilities like `diff`).

Back to the testing problem

- How would you use I/O redirection to speed up the testing process?
 - Write a test generator to generate input files (Use output redirection).
 - Run your program with I/O redirection producing output files.
 - Write a program to compare the generated outputs against a set of expected outputs (Can use existing Linux utilities like `diff`).
- The whole process can be automated with a Bash script (next week)

I/O Redirection

- You can *redirect* standard input, standard output, and standard error to files using I/O redirection
- Use `<` or `0<` to redirect `stdin`
- Use `>` or `1>` to redirect `stdout`
- Use `2>` to redirect `stderr`

I/O Redirection

Notes:

- If the output file already exists for `stdout` and `stderr` redirection, it will be overwritten
 - Use `>>` instead of `>` to *append* to the end of an existing file
- If the output file does not exist, it is created first
- If the standard input is redirected to a non-existing file, this results in an error.
- Most Linux commands are designed to be able to process input from the standard input

echo

```
echo [-n | -e] string
```

Displays the string. It is the *print* statement for the shell.

- `-n` does not print and end of line at the end.
- `-e` turns on recognition of escape sequences (useful for scripting)
 - `\b` backspace
 - `\n` newline
 - `\t` horizontal tab
 - `\v` vertical tab
 - see `man` for more

cat

```
cat [-n] file ...
```

Concatenates each file and displays their content. If only one file is specified, just prints its contents.

- `-n` prints line numbers.
- Can use it with output redirection to merge multiple files to a single file


/dev/null

- If we don't care about a command's output, we can send it to /dev/null
- Output/files sent to /dev/null are deleted forever (recovery not possible)
 - Question: how can you send a file to /dev/null?
- Possible use: If we don't want error messages to clutter output, we can redirect standard error to /dev/null.

Using I/O redirection to connect multiple programs

- If a task requires multiple programs to work on an input in sequence, we can have them communicate with each other using I/O redirection.
- Example: spellchecking a document (See: Kernighan's Unix Demo)
 - Step 1: extract words
 - Step 2: remove punctuation
 - Step 3: convert letters to lowercase
 - Step 4: remove duplicates
 - Step 5: sort the words
 - Step 6: compare against a dictionary

Using I/O redirection to connect multiple programs



Would need a lot of temporary files in between. We will learn a better way to do this with pipes next week.