

210 Systems Programming

The Linux File System, System Administration: Process and User Management

Fall 2025

Week 5

Overview

- Reminder: File permissions
- File system organization
 - inode
 - Symbolic and Hard links
- Becoming a super user
- Querying and manipulating processes

Reminder: Permission Syntax

- Three permission types: read, write, execute (rwx) for three classes: owner, group, others (ugo) \Rightarrow 9 permission flags:
 - $r_u w_u x_u r_g w_g x_g r_o w_o x_o$
- Examples:
 - user has rwx, group has r, others have r \Rightarrow 111 100 100 or `rwxr--r--`
 - user has rw, group has rw, others have none \Rightarrow 110 110 000 or `rw-rw----`

Reminder: Changing permissions with `chmod`

- Can use the *symbolic* or the *numeric* mode
 - The format of a symbolic mode is `[ugoa...][[-+=][perms...]...]`.
 - Examples:
 - `chmod +x foo`
 - `chmod u+rx,g+x,o-rwx foo`
 - The numeric mode use octal digits (0-7) to represent the 3-bit permissions for each ugo category.
 - Examples:
 - `chmod 754 foo`
 - `chmod 600 foo`
- Only the owner of the file/directory or a superuser can change the permissions of that file/directory.

Changing the owner and the group of a file/directory

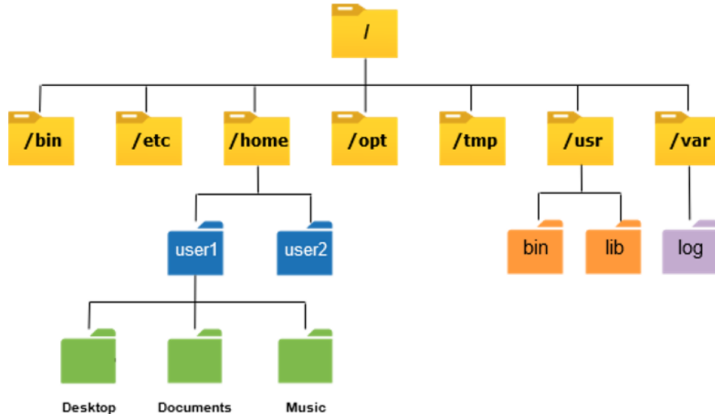
- Only a privileged user can do this. (will learn how to elevate to a *superuser* in the next lecture)

```
chown [OPTION]... [OWNER] [:[GROUP]] FILE...
```

- **C**hange **owner** and group to OWNER and GROUP for each FILE
- Can use the -R option to recursively change the ownership info for the entire content of a directory.

Reminder: Linux Directory Structure

- Files are organized as a set of nested directories



Let's inspect a directory's contents

```
-rw-rw-r-- 1 tolgacan tolgacan      936 Aug 31 2022 log.cpp
-rwxrwxr-x 1 tolgacan tolgacan    17320 Aug 29 2022 myProgram
-rw-rw-r-- 1 tolgacan tolgacan      90 Aug 23 2023 myProgram_backup.cpp
-rw-rw-r-- 1 tolgacan tolgacan      90 Aug 29 2022 myProgram.cpp
-rw-rw-r-- 1 tolgacan tolgacan       0 Jul  9 13:04 newFile
-rw-rw-r-- 1 tolgacan tolgacan     981 Aug 31 2022 orig_log.cpp
drwxrwxr-x 6 tolgacan tolgacan      11 Sep  9 2022 project0-setup-tolgacan
drwxrwxr-x 7 tolgacan tolgacan      17 Sep 23 2022 project1-stack-tolgacan
```

- Consider the file `newFile` in the above output
 - Where is the file name stored?
 - Its size is 0 bytes, so the file name cannot be stored within the file itself.
 - Similarly, where do we store modification date, owner, group, etc.?
 - This additional information about the file is called its *metadata*.
 - In the Linux file system, the metadata is stored separate from the content of the file.

The Unix/Linux File System

A well-formed file system is composed of three main components:

The file

The actual data
blocks containing the
files primary data

inode

A structure assigned
to each file to store
its metadata

Directory structure

A table to store file
names and inode
numbers

Let's look at each of these components in detail and learn how they are all connected.

The File

- A file is a collection of data blocks, on a storage device, containing the data of interest
- For example, if we create a file named `foo` with the following command:
 - `$ echo "Hello World!" > foo`
- The file `foo` is a group of 13 bytes (including the end of line character)
- These blocks do not store the file name `foo`
- This raises a new question:
 - How do we know where these blocks are?

inode

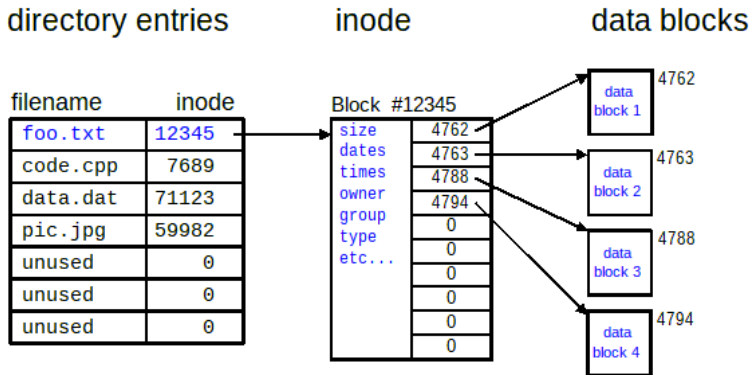
- Each file has an inode to store its metadata and locate the file's contents
 - A regular file's inode contains a pointer to its data blocks
 - A directory's inode contains a pointer to its *directory structure*
- inodes are stored in a table, which is allocated when the filesystem is first created
 - Which means there are a fixed number of inodes in a file system. How is this number determined?
 - In Linux ext4, a default ratio of 1 inode per 16KBs of total capacity is used.
 - Could this be a problem?
- Use the `ls` option `-li` to get the inode number information of files/directories.
- Use `df -li` to get usage information on inodes in a file system.

Directory Structure

- Two pieces of metadata are not stored in the inode:
 - Filename and inode number
- The directory structure represents the contents of each directory as a table of the contained file/directory names and their corresponding inode numbers
- To inspect a directory (instead of listing its contents), use `ls -d`

| Filename | inode # |
|----------|----------|
| a.txt | 33534535 |
| b.txt | 34545221 |
| dir1 | 45455455 |

Three main components together:



From: <https://azrael.digipen.edu/~mmead/www/Courses/CS180/FileSystems-1.html>

Linking: Symbolic Links

- Symbolic links:
 - Same as shortcuts on Windows.
 - They are small (independent and separate) files that contain the *path* information of a target file.
 - When listed inside a directory, the `l` flag at the very first column of a long format listing identifies them as *links*.
 - Also, the target file path is listed along with the symbolic link's file name.
 - Deleting the target file does not remove the symbolic links pointing to that file
 - Create symbolic links using: `ln -s [target file] [link name]`
 - Symbolic links are also called *soft* links

Linking: Hard Links

- Hard links:
 - They are independent entries in the same or different directory structures that point to the **same inode**. In other words, they are aliases to the same single file.
 - You are already familiar with aliases to a directory name: `.`, `..`, and `dirname`,
 - The system keeps track of number of hard links to a file, so that the data blocks are deleted from the disk only after the last remaining link to that inode is deleted.
 - Create hard links using: `ln [target file] [hard link name]`
 - Experiment with hard links and symbolic links and inspect their inode numbers with `ls -li`

The Superuser

- Administers the OS.
 - Is not affected by file permissions.
 - Can move, delete any file that belongs to anyone.
 - Installs, updates applications.
 - Monitors system, network log files.
- Used to be called the *root* user.
- Today, instead of logging as the root user each time, regular users can be given *superuser* privileges by listing them in a special file named: `sudoers`.

Becoming the *superuser*

- While its use is not advised, on some systems you can log into root with `su`.
- This command can be very dangerous, as you will retain the superuser privileges until you exit the *root* shell with the `exit` command.
- Most distributions (i.e. Ubuntu) disable logging into root with `su` altogether.
- Instead, using the `sudo` command to execute **a** command **as** the superuser is recommended.

sudo

```
sudo [OPTIONS] [CMD]
```

Super user do. Run *CMD* as *root*.

- Allows execution of commands with root's privileges, but does not log you in as root.
- Can also be used to execute commands as a different user with the `-u USER` option.
- sudo is generally safe, but we should always be cautious using it.
- Only users in the `/etc/sudoers` may run this command.
- **Do not try it on isengard.** It will be considered as an attack by the system admins.

Installing Programs

- Generally, package managers are used to install programs in an operating systems.
 - On Ubuntu, Debian, and Kali Linux, **apt** is the main package manager.
 - Use, `sudo apt-get update` to update the repository information.
 - Use, `sudo apt-get install package_name` to install a package.
 - On Arch-based distributions, including Arch and Manjaro, **pacman** is used as the package manager.

See: https://en.wikipedia.org/wiki/List_of_software_package_management_systems for more information.

User management

- `who`
 - Prints information about all users who are currently logged in
 - The command `w` is similar and prints more information
- `id [USERNAME]`
 - Prints the user and group information for the specified `[USERNAME]`
- Users can be added using the `useradd` command and can be removed using the `userdel` command
 - Must have superuser privileges

What is a *process*?

- A *process* is a *running* copy of a program along with its allocated resources such as, memory, file handles, network sockets, etc.
- There can be multiple different processes of a program.
 - Example: vim is a program installed on *isengard*. When different users edit their own files simultaneously, there will be multiple and separate *vim* processes with their own process ids, memory footprints, etc.
- Processes are dynamic, i.e., there are various events in their lifetimes:
 - Created, Paused, Run, Killed, Completed

How to get a list of current processes?

- Use the `ps` command to get a list of processes.
 - `ps -ef` will list all the processes in full-format.
 - Check `man ps` for different options
 - Demo on *isengard*

Other tools for process monitoring

- `top` and `htop` are two interactive programs for monitoring processes and system resources such as CPU and memory.
- Try them out on `isengard` (or on your own Linux)

Foreground and background processes

- When you run a program in your bash shell, by default, it is run as a *foreground* process, meaning that your interaction with the shell is suspended until that program finishes its execution.
- Most programs finish in an instant, so this is not a big deal. However, imagine you have a program that runs for 1-2 hours or maybe for days.
 - You would want to run it as a *background* process to keep it running in the background why you get back to your shell prompt instantly.
 - Use `<program> &` to run a program as a background process.

Sending pause/kill signals to a foreground process

- A running foreground process can be terminated by sending it a SIGINT signal with Ctrl-C.
 - We will learn more about signals when we talk about processes in Module 3.
- A running foreground process can be paused by sending it a SIGTSTP signal with Ctrl-Z.
 - The most recently paused process can resume execution in the foreground with a fg command and in the background with a bg command.

Using the `kill` command to terminate/pause processes

- If you want to terminate or pause a process you own that is running in the background, you can use the `kill` command, to send it `SIGTERM` or `SIGTSTP` signals. Examples:
 - `kill -TSTP pid` to pause (sends `SIGTSTP` signal)
 - `kill -CONT pid` to resume execution
 - `kill pid` to terminate (sends `SIGTERM` signal)
- Alternatively, you can use the `pkill` command to send signals processes based on their name instead of their process ids.

Why is there a *parent* process?

- When you list the processes using `ps -ef`, you see that each process has a PPID (Parent Process ID) listed along with other information.
- What does this tell you about how processes are created?

Process creation

- We will talk about this in more detail in Module 3; but, shortly, each process is created by a parent process.
 - The parent process uses the `fork()` system call to create a clone of itself.
 - The clone is identical to the current process except its process id and parent process id
 - The new process then loads and executes the desired program

fork() example

- When `fork()` is successful, it returns two different outputs to the child and the parent process.
 - Child PID is returned to the parent
 - 0 is returned to the child
- When `fork()` fails, no child process is created and -1 is returned to the parent.

fork() example

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();

    if (rc < 0) {
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) { // child (new process)
        printf("hello, I am child of %d (my pid:%d)\n", (int) getppid(), (int) getpid());
    } else { // parent goes down this path (original process)
        printf("hello, I am parent of %d (my pid:%d)\n", rc, (int) getpid());
    }
    return 0;
}
```