

210 Systems Programming

Pipes, text processing, grep

Fall 2025

Week 3

Overview

- Piping
- `wc`, `head`, `tail`, `tr`, `sort`, and `uniq` commands
- The `grep` command
 -
 - Standard `grep`
 - Extended regular expressions `egrep`

A Better Way to connect I/O: Pipes

A mechanism to channel standard input and standard output between programs.

- Syntax: `s1 | s2 | s3 | ...`
- Chain together multiple commands/programs
- Data flows left to right
- The first program gets input from the standard input, the last program writes to standard output
- The standard output of the first program becomes the standard input of the second program
- ...

Text Processing with Pipes

- There are a suite of commands made especially for piping
- We'll go through a few in this lecture and a very important one (grep) next lecture

head and tail

```
head [-n number] file
```

Prints the first number lines of a file. By default the first 10 lines are displayed.

```
tail [-n number] file
```

Prints the last number lines of a file. By default the last 10 lines are displayed.

- Question: How can you display lines 56-60 of a file that has 1000 lines using pipes?

head and tail

```
head [-n number] file
```

Prints the first number lines of a file. By default the first 10 lines are displayed.

```
tail [-n number] file
```

Prints the last number lines of a file. By default the last 10 lines are displayed.

- Question: How can you display lines 56-60 of a file that has 1000 lines using pipes?
- Answer: `head -n 60 file.txt | tail -n 5`

sort

```
sort [-r] [-k field1[,field2]] [file ...]
```

Sorts a file (or files) line by line.

- `-r` sorts in reverse order
- Use `-k` to sort by a specific column (default: first column)
- Note that the `file` input is optional. Why?
- Sorting order:
 - Special characters, numbers, letters (lowercase before uppercase of same letter)

tee

```
tee [-a] file
```

The tee utility copies standard input to standard output, making a copy in file.

- -a Append the output to the file rather than overwriting it.

- Example:

- `echo "Hello" | tee greetings.txt`
Hello

uniq

```
uniq [-c|-d|-u] file
```

Check for **unique** lines in file.

- Behavior is greatly dependant on given options.
- By default, file is printed to stdout with duplicates removed
- IMPORTANT: Input must be sorted! Duplicates are only found if adjacent
- -c count occurrences of each line
- -d output just the duplicate lines
- -u output just unique lines

tr

```
tr [-C] string1 string2 | tr -d string1
```

The `tr` utility either substitutes the characters in `string1` to characters in `string2` or deletes characters in `string1` in standard input and outputs the result in standard output.

- Note: does not take file input
- Use `-C` to complement the characters in `string1`

- Example:

- `cat file.txt | tr "[a-z]" "[A-Z]" > uppercase.txt`

WC

```
wc [-l | -w | -c | -m] file ...
```

Count the number of lines, words, bytes, or characters in file

- Use `-l` to count lines
- Use `-w` to count words
- Use `-c` to count bytes
 - Same as `-m` if multibyte characters are not present
- Use `-m` to count characters

grep History



- Ken Thompson, one of the inventors of Unix, was helping a fellow coworker do some textual analysis on The Federalist Papers.
- Thompson written his own program that allowed text searching by using regular expressions.
- He named this tool “**G**lobal **R**egular **E**xpression **P**rint”, or simply grep

grep History



- Thompson's boss, Doug McIlroy, approached him about the need for a text searching utility.
- Thompson promised to work on it overnight, but really only spent about an hour fixing bugs, since he'd already written Grep and had been using it privately.
- He presented it to McIlroy the next day. And the rest is history.

Simple grep

```
grep [-i|-c|-l|-n|-v|-o|-R] pattern file ...
```

Search for pattern in each file and print matched lines

- -i ignore-case
- -c return total match count (**of lines**) instead of line contents
- -l return names of matched files, instead of line contents
- -n show line numbers
- -v return lines which do not match pattern
- -o print only the matching parts on separate lines
- -R read the files in directories, recursively



Using grep

- Learning how to use grep effectively boils down to learning the language to define patterns: *regular expressions*
- grep manual at gnu.org
 - <https://www.gnu.org/software/grep/manual/>
 - A 43-page document!

Extended Regular Expression Syntax

Special characters: `.?*\{ | () [\ ^ $`

All the other characters are ordinary characters.

`?*\{` are repetition operators. The ones beginning with `{` are called *interval expressions*:

<i>Symbol</i>	<i>Meaning</i>
<code>.</code>	Matches any single character
<code>?</code>	Matches preceding item 0 or once
<code>*</code>	Matches preceding item 0 or more times
<code>+</code>	Matches preceding item 1 or more times

Extended Regular Expression Syntax

Repetition operators continued

<i>Symbol</i>	<i>Meaning</i>
$\{n\}$	Matches preceding item exactly n times
$\{n, \}$	Matches preceding item n or more times
$\{, m\}$	Matches preceding item at most m times
$\{n, m\}$	Matches preceding item at least n times, but not more than m times.

Extended Regular Expression Syntax

Additional Notes

- The empty regular expression matches the empty string.
- Two regular expressions may be concatenated; the resulting regular expression matches any string formed by concatenating two substrings that respectively match the concatenated expressions.
- Two regular expressions may be joined by '|'. Either of the two expressions, which are called *alternatives*, are matched.
- Repetition takes precedence over concatenation, which in turn takes precedence over alternation.
- A whole expression may be enclosed in parentheses to override precedence rules. An unmatched ')' matches just itself.

RE Syntax

Bracket Expressions

- A *bracket expression* is a list of characters enclosed by '[' and ']'.
- It matches any single character in that list.
- If the first character of the list is the caret ^, then it matches any character **not** in the list.
- Examples:
 - [0123456789] matches any single digit.
 - [^()] matches any single character that is not opening or closing parenthesis.
- Special characters lose their special meaning inside bracket expressions.

RE Syntax

Range Expressions within Brackets

- Within a bracket expression, a *range* expression consists of two characters separated by a hyphen.
- It matches any single character that sorts between the two characters, inclusive.
- In the default C locale, the sorting sequence is the native character order; for example, `'[a-d]'` is equivalent to `'[abcd]'`.
- To obtain the traditional interpretation of bracket expressions, you can use the 'C' locale by setting the `LC_ALL` environment variable to the value 'C'.

RE Syntax

Character Classes

<i>Symbol</i>	<i>Meaning</i>
<code>[:alnum:]</code>	Alphanumeric characters. Same as <code>[0-9A-Za-z]</code>
<code>[:alpha:]</code>	Alphabetic characters. Same as <code>[A-Za-z]</code>
<code>[:lower:]</code>	Lowercase letters
<code>[:upper:]</code>	Uppercase letters
<code>[:digit:]</code>	Digits
<code>[:xdigit:]</code>	Hexadecimal digits
<code>[:blank:]</code>	Blank characters: space and tab
<code>[:punct:]</code>	Punctuation characters

RE Syntax

Special Backslash Expressions

The ‘\’ character, when followed by certain ordinary characters, takes a special meaning:

<i>Symbol</i>	<i>Meaning</i>
<code>\b</code>	Matches the empty string at the edge of a word
<code>\B</code>	Match the empty string provided it's not at the edge of a word
<code>\<</code>	Matches the empty string at the beginning of a word
<code>\></code>	Matches the empty string at the end of a word

Example:

- ‘`\brat\b`’ matches the separate word ‘rat’, ‘`\Brat\B`’ matches ‘crate’ but not ‘furry rat’.

RE Syntax

Anchoring

- The caret '^' and the dollar sign '\$' are special characters that respectively match the empty string at the beginning and end of a line.
- They are termed anchors, since they force the match to be *anchored* to beginning or end of a line, respectively.

RE Syntax

Differences of the Basic RE

Basic RE has functionally the same power as the Extended RE. However, Basic regular expressions differ from extended regular expressions in the following ways:

- The characters `?`, `+`, `{`, `|`, `(`, and `)` lose their special meaning; instead use the backslashed versions.
- If an unescaped `^` (or `$`) appears neither first (last), nor directly after (before) `\(` or `\|`, it is treated like an ordinary character and is not an anchor.
- If an unescaped `*` appears first, or appears directly after `\(` or `\|` or anchoring `^`, it is treated like an ordinary character and is not a repetition operator.

grep examples

- `grep ^alias .bashrc`
 - Find lines in your `.bashrc` file that **start with** `alias`
- `grep -n ^$ *.c`
 - Find empty lines in all the `c` files in the current directory and report their line numbers
- `grep '\bse[et]\b' .bashrc`
 - Find occurrences of `see` or `set` that appear as whole words in `.bashrc`
- `grep -o -n -E '(0|[1-9][0-9]*)([eE][+-]?[0-9]+)?' *`
 - Find integers (possibly in scientific notation, as well) and print integers only with line numbers
- See more at: <https://phoenixnap.com/kb/grep-regex>